



a place of mind

THE UNIVERSITY OF BRITISH COLUMBIA

# Dynamic memory

Geoff's self-checklist:

- ☐ There is a new iClicker Assignment
- ☐ Record lecture

# Announcements

- Midterm scheduling
  - Piazza poll reopened – please respond!
- Lab 2 starts next week!

# iClicker assignment

## Question 1

- What happens when we run the code below?

```
void plusfive(int x) {  
    x = x + 5;  
}  
  
int main(void) {  
    int y = 10;  
    plusfive(y);  
    printf("%d\n", y);  
    return 0;  
}
```

Assume that the necessary libraries have been #include-d

- A. 10 is printed
- B. 15 is printed
- C. Something else is printed
- D. There will be a runtime error
- E. The code does not compile

# iClicker assignment

## Question 2

- What happens when we run the code below?

```
void plusfive(int* x) {  
    *x = *x + 5;  
}  
  
int main(void) {  
    int y = 10;  
    plusfive(&y);  
    printf("%d\n", y);  
    return 0;  
}
```

Assume that the necessary libraries have been #include-d

- A. 10 is printed
- B. 15 is printed
- C. Something else is printed
- D. There will be a runtime error
- E. The code does not compile

# iClicker assignment

## Question 3

- What would be the result of executing this code?

Assume that the necessary libraries have been #include-d

- A. 12 is printed
- B. 16 is printed
- C. 17 is printed
- D. 20 is printed
- E. Nothing is printed due to some error

```
int main(void) {  
    int arr[] = {3, 4, 5, 6, 7, 8};  
    int* q;  
    int i;  
    int sum = 0;  
  
    for (q = &arr[1]; q < &arr[4]; q++) {  
        *q = *q + 1;  
    }  
  
    for (i = 0; i < 6; i += 2) {  
        sum += arr[i];  
    }  
  
    printf("%d", sum);  
  
    return 0;  
}
```

# Dynamic memory

- Arrays declared as local variables must have a known size at compile time
  - but sometimes we don't know how much space we need until runtime
- Suppose we expect users to only need to store up to 1000 values, so we hard-code "1000" as an array size
  - What if user needs more? Change code and recompile
  - What if user only needs 5? Wastes memory
- If the value 1000 is hard-coded, this is hard to find and change
  - especially if used in multiple locations
- If hardcoded as a symbolic constant, still cannot change without recompiling

# Memory management in C

- We have already seen how locally-declared variables are placed on the function call stack
  - allocation and release are managed automatically
- The available stack space is extremely limited
  - placing many large variables or data structures on the stack can lead to stack overflow
- Stack variables only exist as long as the function that declared them is running

# Dynamic memory allocation

- At run-time, we can request extra space on-the-fly, from the *memory heap*
- Request memory from the heap – "allocation"
- Return allocated memory to the heap (when we no longer need it) – "deallocation"
- Unlike stack memory, items allocated on the heap must be explicitly freed by the programmer



# Dynamic memory allocation

- Function `malloc` returns a pointer to a memory block of at least `size` bytes:

```
ptr = (cast-type*) malloc(byte-size);
```


- Function `free` returns the memory block (previously allocated with `malloc`) and pointed to by `ptr` to the memory heap:

```
free(ptr);
```

- The system knows how many bytes need to be freed, provided we supply the correct address `ptr`

# Heap example

```
int main() {  
    int a;  
    int *p = (int*) malloc(sizeof(int));  
    *p = 10;  
}
```



- If there is no free memory left on the heap, **malloc** will return a null pointer
- Note: **malloc** only allocates the space but does not initialize the contents.
  - Use **calloc** to allocate and clear the space to binary zeros

# Allocating dynamic arrays

- Suppose we want to allocate space for exactly 10 integers in an array

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int* i;
    i = (int*) malloc(10*sizeof(int));
    if (i == NULL) {
        printf("Error: can't get memory...\n");
        exit(1); // terminate processing
    }

    i[0] = 3; // equivalent: *(i+0) = 3;
    i[1] = 16; // *(i+1) = 16;
    printf("%d", *i);
    ...
}
```

# Allocating dynamic arrays

From user input, variable array size

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int employees, index;
    double* wages;
    printf("Number of employees? ");
    scanf("%d", &employees);

    wages = (double*) malloc(employees * sizeof(double))
    if (!wages) { // equivalent: if (wages == NULL)
        printf("Error: can't get memory...\n");
    }

    printf("Everything is OK\n");
    ...
}
```

See `dma_examples.c`

# Dangling pointers

- When we are done with an allocated object, we free it so that the system can reclaim (and later reuse) the memory

```
int main() {  
    int* i = (int*) malloc(sizeof(int));  
    *i = 5;  
    free(i);  
  
    printf("%d", *i);  
}
```

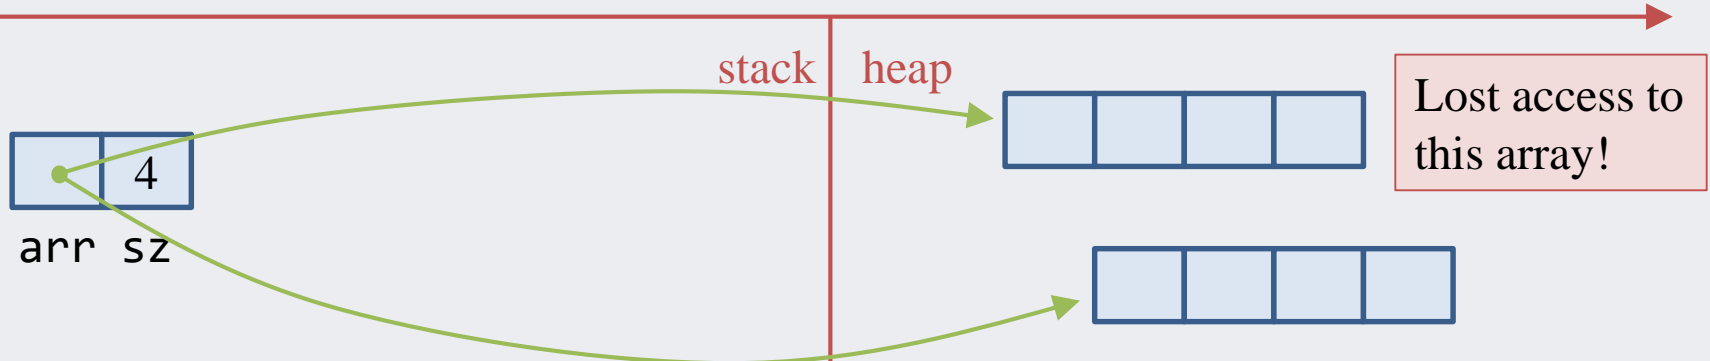
The space is marked as free, but the value remains until it is overwritten

- If the pointer continues to refer to the deallocated memory, it will behave unpredictably when dereferenced (and the memory is reallocated) – a **dangling pointer**
  - Leads to bugs that can be subtle and brutally difficult to find
  - So, set the pointer to **NULL** after freeing `i = NULL;`

# Memory leaks

- If you lose access to allocated space (e.g. by reassigning a pointer), that space can no longer be referenced, or freed
  - And remains marked as allocated for the lifetime of the program

```
int* arr;  
int sz = 4;  
arr = (int*) malloc(sz*sizeof(int));  
arr[2] = 5;  
arr = (int*) malloc(sz*sizeof(int));  
arr[2] = 7;
```



See `memory_leak.c`

# Exercise


- What is printed to the screen?
  - Also clearly identify memory leaks and dangling pointers

```
int w;  
int z;  
int* t = (int*) malloc(sizeof(int));  
int* y = (int*) malloc(sizeof(int));  
int* x = (int*) malloc(sizeof(int));  
  
*x = 3;  
*y = 5;  
z = *x + *y;  
w = *y;  
*x = z;  
free(x);  
*t = 2;  
y = &z;  
x = y;  
free(t);  
printf("*x=%d, *y=%d, z=%d, w=%d\n", *x, *y, z, w);
```

# Dynamic allocation of a 2D array

See dma\_2d.c for details

```
int dim_row = 3;  
int dim_col = 5;  
int** myarray; // pointer to a pointer
```





# Stack memory vs heap memory

- Stack
  - fast access
  - allocation/deallocation and space automatically managed
  - memory will not become fragmented
  - local variables only
  - limit on stack size (OS-dependent)
  - variables cannot be resized
- Heap
  - variables accessible outside declaration scope
  - no (practical) limit on memory size
  - variables can be resized
  - (relatively) slower access
  - no guaranteed efficient use of space
  - memory management is programmer's responsibility

# Readings for this lesson

- Thareja
  - Appendices A, B, E
- Next class:
  - Thareja, Chapters 4 – 5